

Intelligent Systems | HS2025 | SW05

Example of simple linear regression with python, SW05 lecture

Eugen Rodel, 15/10/2025

Forecasting Students Grade

1. Import Libraries

```
In [14]: import pandas as pd
import numpy as np
import matplotlib.pyplot as plt
import seaborn as sns
from sklearn.model_selection import train_test_split
from sklearn.linear_model import LinearRegression
from sklearn.metrics import mean_squared_error, r2_score

# Set Seaborn style for plots
sns.set_style('whitegrid')
```

2. Load Data


```
In [15]: # Load the uploaded data from the CSV file
data = pd.read_csv('study_hours_grades_50.csv')

# Display the first few rows of the DataFrame to understand the structure
data.head()
```

```
Out[15]:   Study Hours  Grade (%)
0            0.00      52.5
1            0.20      50.3
2            0.41      55.3
3            0.61      60.7
4            0.82      52.9
```

3. Visualize Data

```
In [16]: # Scatter plot to visualize the relationship between Gas Pumps and Sales per Year
sns.scatterplot(x='Study Hours', y='Grade (%)', data=data)
plt.xlabel('Study Hours')
plt.ylabel('Grade (%)')
plt.title('Scatter Plot of Grade (%) vs Study Hours')
plt.show()
```



4. Prepare Data

```
In [17]: # Separate the features/independent variables (Study Hours) and the target/dependent variable
X = data[['Study Hours']].values
y = data['Grade (%)'].values
```



```
In [18]: # Split the data into training and testing sets
X_train, X_test, y_train, y_test = train_test_split(X, y, test_size=0.2, random_state=0)
```

5. Train The Model

```
In [19]: # Initialize the Linear Regression model
model = LinearRegression()

# Fit (or Learn) the model to the training data
model.fit(X_train, y_train)
```



```
Out[19]: ▾ LinearRegression ⓘ ⓘ
LinearRegression()
```

6. Display Slope and Intercept

```
In [20]: # Display the slope (coefficient) and intercept of the trained model
print(f'Coefficient (Slope): {model.coef_[0]:.2f}')
print(f'Intercept: {model.intercept_:.2f}')
```

Coefficient (Slope): 4.67
 Intercept: 50.27

7. Make Predictions

```
In [21]: # Predict the target variable for the test set  
y_pred = model.predict(X_test)
```

8. Evaluate The Model

```
In [22]: print(f'Test Data (How many study hours):\n{X_test}\n')  
print(f'Predictions from trained model (Predicted grade (%)):\n {y_pred}')
```

Test Data (How many study hours):

```
[[5.71]  
[2.24]  
[2.04]  
[8.37]  
[0.41]  
[5.51]  
[7.76]  
[6.33]  
[4.49]  
[0.82]]
```

Predictions from trained model (Predicted grade (%)):

```
[76.96039036 60.73945819 59.80453415 89.39488015 52.18490319 76.02546632  
86.54336182 79.8586549 71.25735369 54.10149748]
```

9. Calculate Mean Squared Error and R-squared

```
In [23]: # Calculate Mean Squared Error and R-squared  
mse = mean_squared_error(y_test, y_pred)  
r2 = r2_score(y_test, y_pred)  
  
# Output the evaluation metrics  
print(f'Mean Squared Error: {mse:.2f}')  
print(f'R-squared: {r2:.2f}')
```


Mean Squared Error: 18.64

R-squared: 0.91

10. Visualize the Linear Regression Line

```
In [24]: # Create a scatter plot for the actual data points  
sns.scatterplot(x=X_test.flatten(), y=y_test, color='red', label='Actual', s=50)  
  
# Create a line plot for the regression line (predicted values)  
sns.lineplot(x=X_test.flatten(), y=y_pred, color='blue', linewidth=2, label='Predicted')  
  
# Label the plot  
plt.xlabel('Study hours')  
plt.ylabel('Grade (%)')  
plt.title('Linear Regression Fit: Grade (%) vs Study hours')  
plt.legend()  
plt.show()
```

Linear Regression Fit: Grade (%) vs Study hours



```
In [25]: # New grade with 2.5 hours of study
new_study = np.array([[2.5]])

# Use the trained model to predict sales for the new gas station
predicted_grade = model.predict(new_study)

# Output the prediction
print(f'Predicted Grade for a 2.5 hours of study: {predicted_grade[0]:.2f} (in %)')

Predicted Grade for a 2.5 hours of study: 61.95 (in %)
```

```
In [ ]:
```