

IN_SYS - SW06 - Multiple Linear Regression - California Housing Data

30/10/2025 - Eugen Rodel - v2

Goals:

- Deepen understanding of Multiple Linear Regression using California Housing Data
- Repeat old topics as importing data from CSV file into Pandas, exploring and visualization of data, train model, make predictions

1. California Housing Dataset Explanation

The **California_Housing.csv** dataset contains information about housing features in California, such as median income, number of rooms, and house prices. Below is a description of each column in the dataset:

1. Median Income:

- Median income of households in the area, measured in tens of thousands of USD.
- This feature is an indicator of the economic status of the neighborhood.

2. Age of House:

- Age of the house in years.
- Represents the average age of houses in the neighborhood.

3. Average Rooms:

- Average number of rooms per household in the area.
- Includes all rooms such as bedrooms, living rooms, kitchens, etc.

4. Average Bedrooms:

- Average number of bedrooms per household in the area.
- Represents the number of bedrooms available in homes, on average.

5. Population:

- Total population of the area.
- Indicates the number of people residing in the neighborhood.

6. Average Occupancy:

- Average occupancy per household in the area.
- Represents how many people live in each household, on average.

7. Latitude:

- Latitude of the location.

- Represents the geographic coordinate for the neighborhood.

8. Longitude:

- Longitude of the location.
- Represents the geographic coordinate for the neighborhood.

9. Price:

- Price of the house in hundreds of thousands of USD.
- This is the **dependent variable** that we are trying to predict in our regression model. It represents the average housing price for the neighborhood.

The goal of this dataset is to predict the **Price** based on various features related to the property and the surrounding area. This dataset provides a practical use case for understanding how different factors, such as house size, age, and income, impact housing prices.

2. Import libraries

Hint: no new libraries are needed, just what we used in the lecture for the advertising Multiple Linear Regression.

```
In [1]: import pandas as pd
import seaborn as sns
import matplotlib.pyplot as plt
from sklearn.model_selection import train_test_split
from sklearn.linear_model import LinearRegression
from sklearn.metrics import mean_squared_error, r2_score

# Set Seaborn style for plots
sns.set_style('whitegrid')
```

3. Load the Dataset

```
In [2]: # Load the 'California_Housing.csv' file into a Pandas DataFrame called 'housing'
#housing_df = pd.read_csv('California_Housing.csv')
housing_df = pd.read_csv('California_Housing.csv')

# Display the first few rows of the DataFrame to verify that the import worked
housing_df.head()
```

Out[2]:

	Median Income	Age of House	Average Rooms	Average Bedrooms	Population	Average Occupancy	Latitude	Longitude	Price
0	8.3252	41.0	6.984127	1.023810	322.0	2.555556	37.88	-122.23	465000
1	8.3014	21.0	6.238137	0.971880	2401.0	2.109842	37.86	-122.22	300000
2	7.2574	52.0	8.288136	1.073446	496.0	2.802260	37.85	-122.24	300000
3	5.6431	52.0	5.817352	1.073059	558.0	2.547945	37.85	-122.25	300000
4	3.8462	52.0	6.281853	1.081081	565.0	2.181467	37.85	-122.25	300000

4. Data Exploration and Visualization

In [3]: `# Display shape of the DataFrame
housing_df.shape`

Out[3]: (20640, 9)

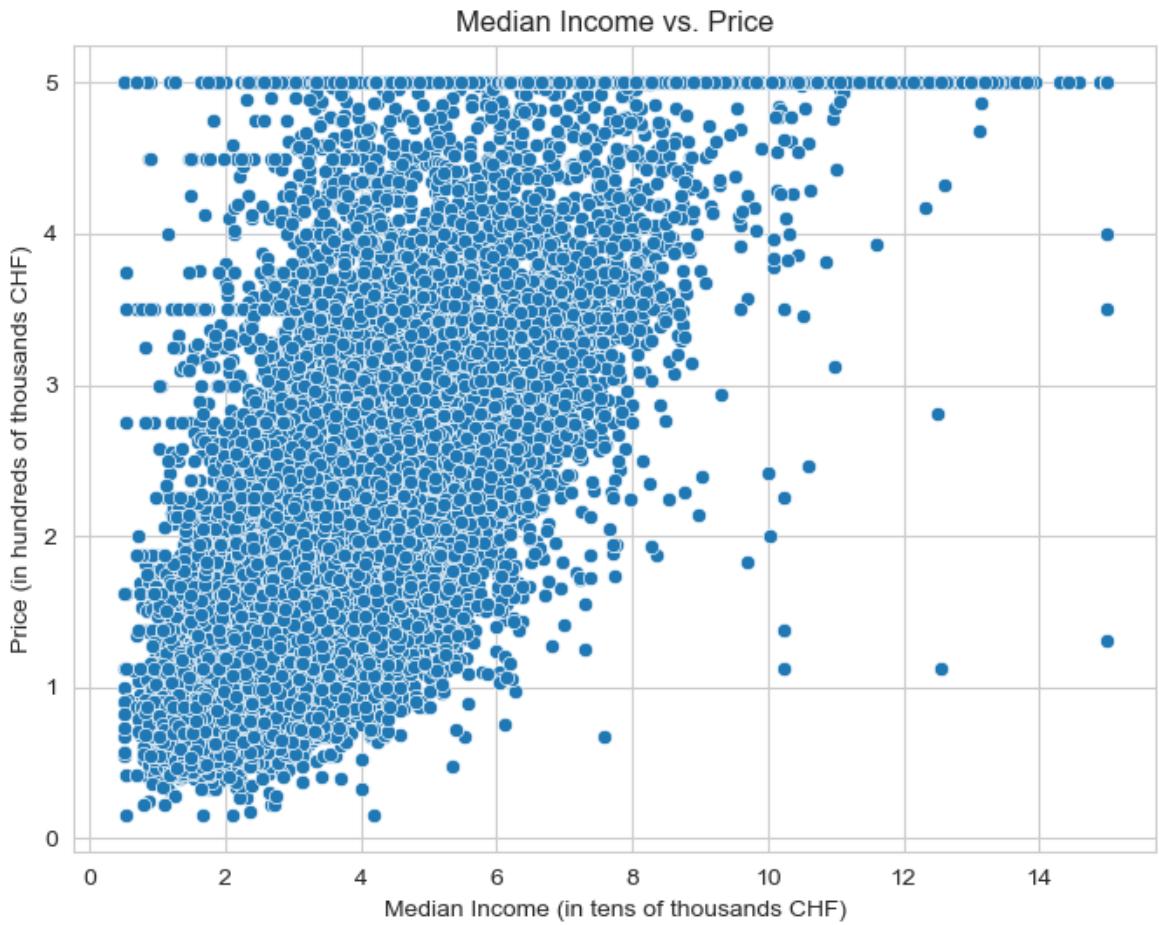
In [4]: `# Display column names
housing_df.columns`

Out[4]: Index(['Median Income', 'Age of House', 'Average Rooms', 'Average Bedrooms', 'Population', 'Average Occupancy', 'Latitude', 'Longitude', 'Price'],
dtype='object')

In [5]: `# Display summary statistics
housing_df.describe()`

	Median Income	Age of House	Average Rooms	Average Bedrooms	Population	Average Occupancy
count	20640.000000	20640.000000	20640.000000	20640.000000	20640.000000	20640.000000
mean	3.870671	28.639486	5.429000	1.096675	1425.476744	3.070
std	1.899822	12.585558	2.474173	0.473911	1132.462122	10.386
min	0.499900	1.000000	0.846154	0.333333	3.000000	0.692
25%	2.563400	18.000000	4.440716	1.006079	787.000000	2.429
50%	3.534800	29.000000	5.229129	1.048780	1166.000000	2.818
75%	4.743250	37.000000	6.052381	1.099526	1725.000000	3.282
max	15.000100	52.000000	141.909091	34.066667	35682.000000	1243.333

In [6]: `# Scatter plot of Median Income vs Price
plt.figure(figsize=(8, 6))
sns.scatterplot(data=housing_df, x='Median Income', y='Price', marker='o')
plt.title('Median Income vs. Price')
plt.xlabel('Median Income (in tens of thousands CHF)')
plt.ylabel('Price (in hundreds of thousands CHF)')
plt.show()`



5. Data "Cutting"

In the preceding scatter plot, we can see that the data has not been cleaned (e.g. where the price is constant and independent of income). A quick and dirty solution is to only use data where the price is < 3.

```
In [7]: # copy the data
housing_cut_df = housing_df.copy()

# filter the data
housing_cut_df = housing_cut_df[housing_cut_df['Price'] < 3]
```

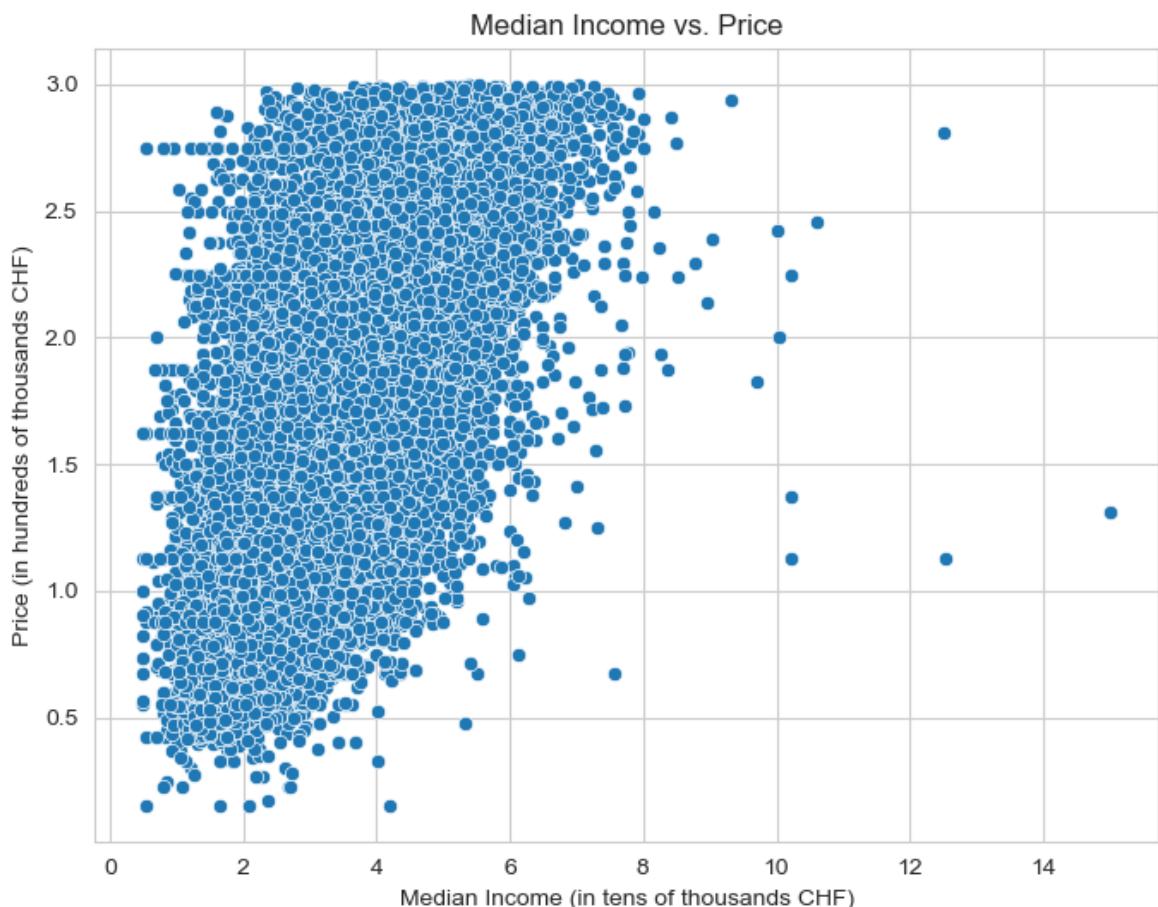
```
In [8]: # describe data
housing_cut_df.describe()
```

Out[8]:

	Median Income	Age of House	Average Rooms	Average Bedrooms	Population	Aver Occupa
count	16774.000000	16774.000000	16774.000000	16774.000000	16774.000000	16774.000000
mean	3.383996	27.979790	5.263933	1.101897	1457.075176	3.153
std	1.312453	12.307397	2.329034	0.477026	1148.226896	10.844
min	0.499900	1.000000	0.846154	0.333333	5.000000	0.692
25%	2.392900	18.000000	4.370574	1.006734	804.000000	2.495
50%	3.221600	28.000000	5.103504	1.050108	1196.000000	2.894
75%	4.208100	36.000000	5.821055	1.102007	1771.000000	3.374
max	15.000100	52.000000	132.533333	34.066667	35682.000000	1243.333

In [9]:

```
# Scatter plot of Median Income vs Price
plt.figure(figsize=(8, 6))
sns.scatterplot(data=housing_cut_df, x='Median Income', y='Price', marker='o')
plt.title('Median Income vs. Price')
plt.xlabel('Median Income (in tens of thousands CHF)')
plt.ylabel('Price (in hundreds of thousands CHF)')
plt.show()
```



In [10]:

```
# Create a heatmap to show correlations between features
plt.figure(figsize=(8, 6)) # Set the plot size in inch: figsize(width, height)
sns.heatmap(housing_cut_df.corr(), annot=True, cmap='coolwarm', linewidths=0.5)
plt.title('Correlation Heatmap')
```

```

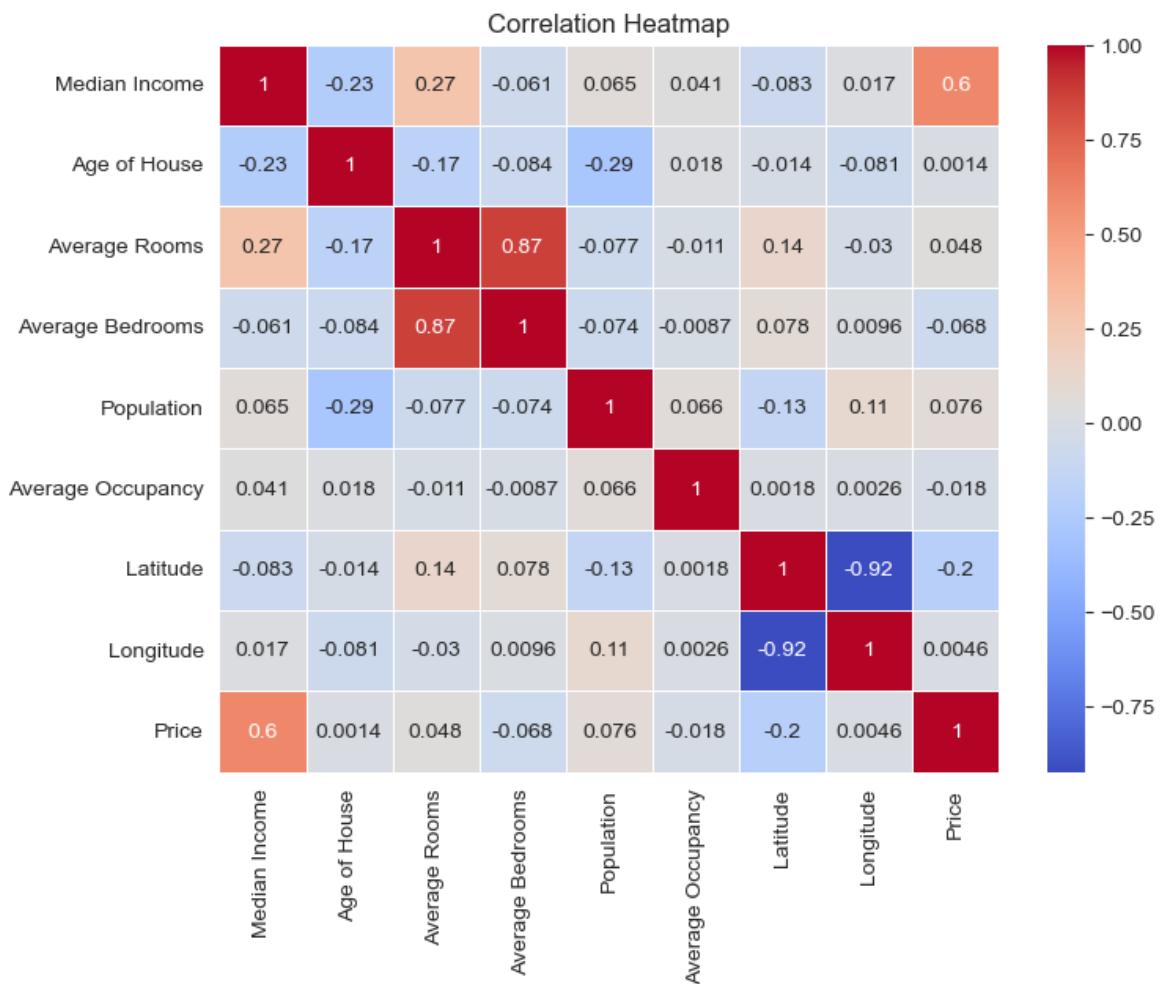
# Box plot for "Median Income"
plt.figure()
sns.boxplot(data=housing_cut_df, y='Median Income')
plt.title('Box Plot of Median Income')
plt.ylabel('Median Income')

# Box plot for "Age of House"
plt.figure()
sns.boxplot(data=housing_cut_df, y='Age of House')
plt.title('Box Plot of Age of House')
plt.ylabel('Age of House')

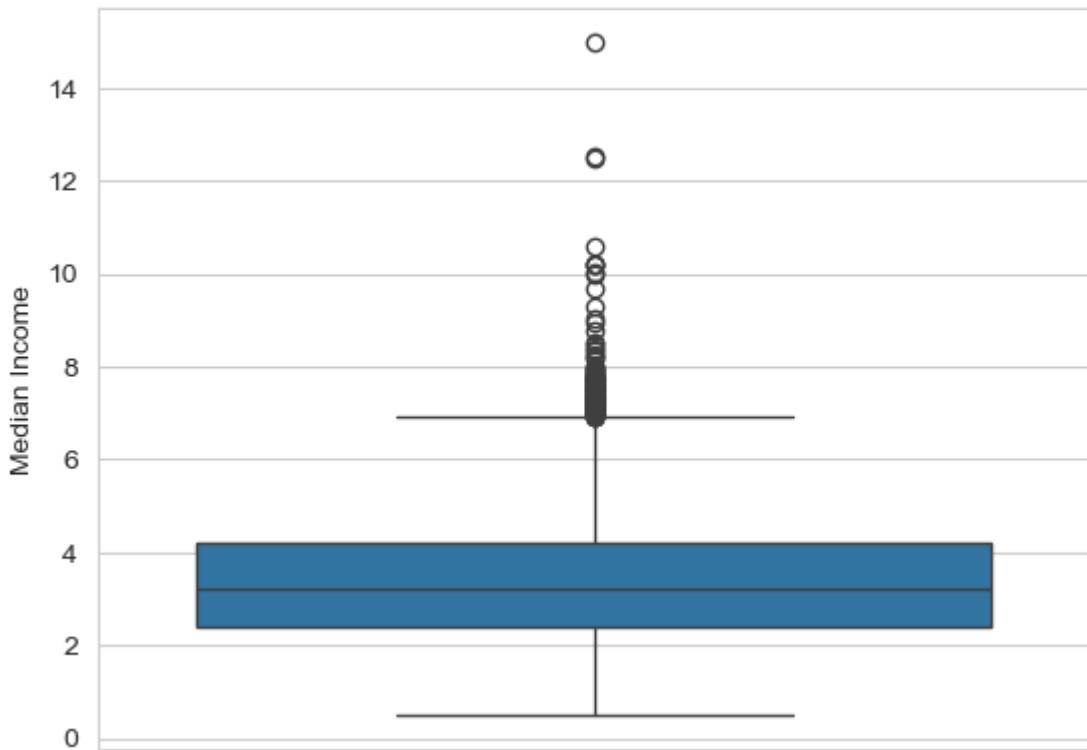
# Box plot for "Price"
plt.figure()
sns.boxplot(data=housing_cut_df, y='Price')
plt.title('Box Plot of Price')
plt.ylabel('Price')

```

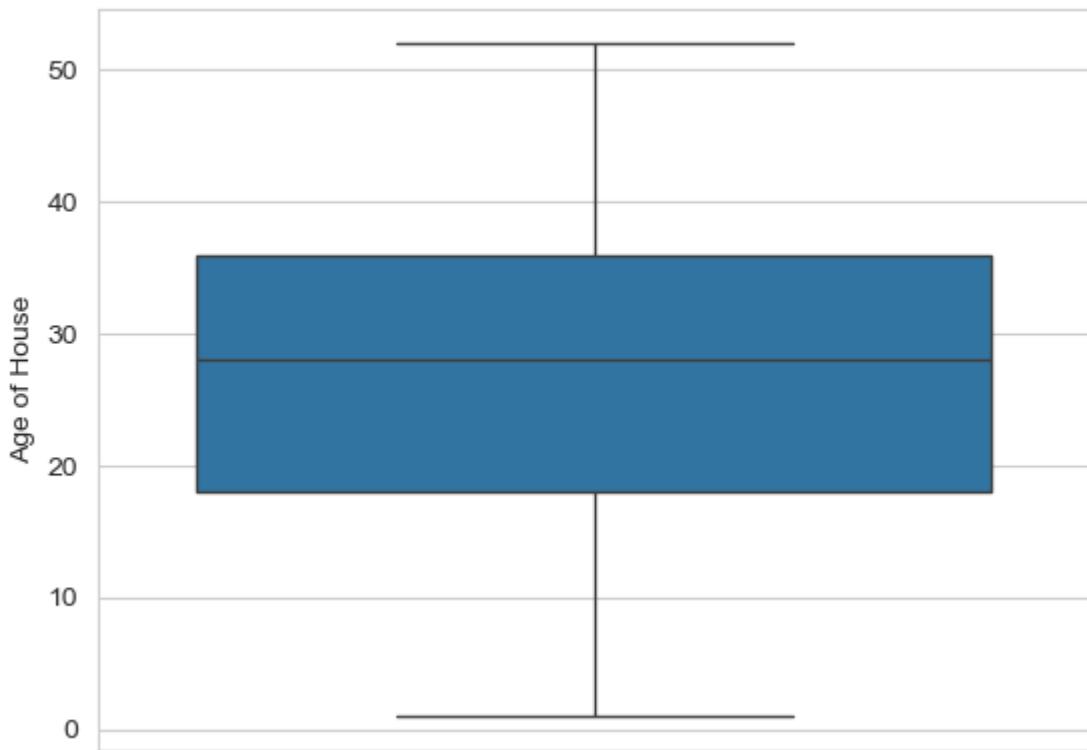
Out[10]: Text(0, 0.5, 'Price')

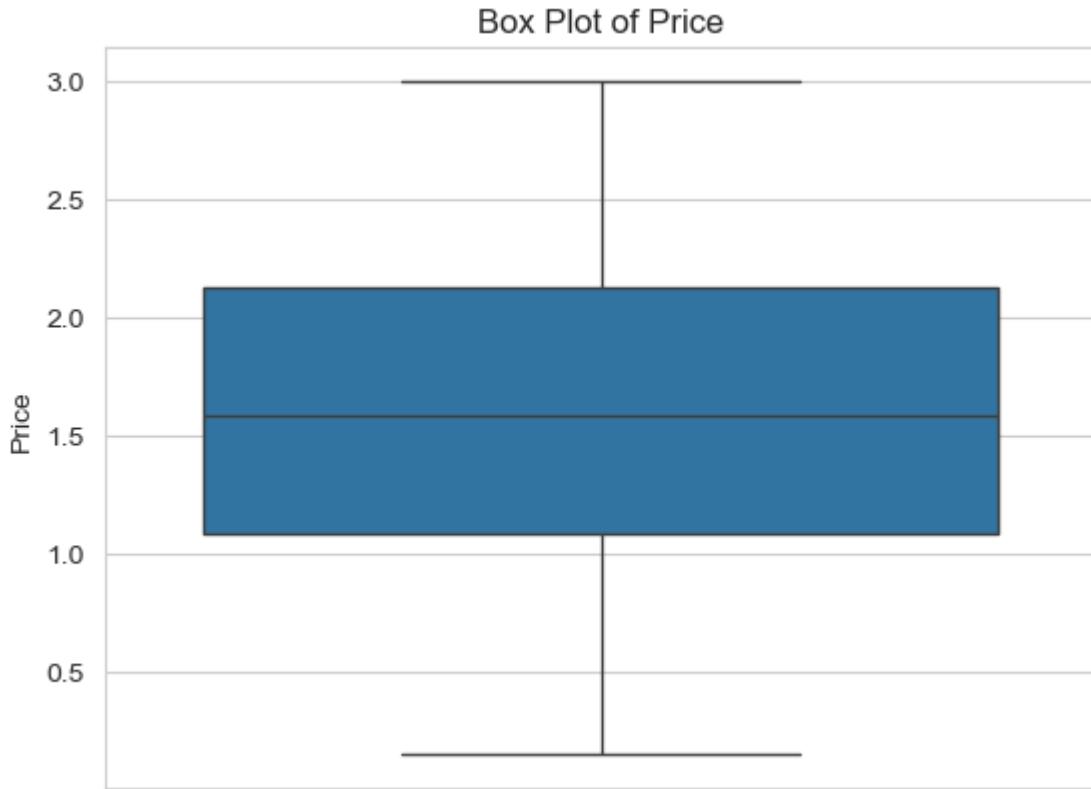


Box Plot of Median Income



Box Plot of Age of House





6. Prepare the Data for Modeling

```
In [11]: # Define independent input features (X) and dependend output target - the sales
X = housing_cut_df[['Median Income', 'Age of House', 'Average Rooms', 'Average B
y = housing_cut_df['Price']

In [12]: # Split the data into training and testing sets (80% training, 20% testing), with
X_train, X_test, y_train, y_test = train_test_split(X, y, test_size=0.2, random_
```

7. Train the Multiple Linear Regression Model

```
In [13]: # Initialize and train the Linear Regression model
model = LinearRegression()
model.fit(X_train, y_train)

Out[13]: ▾ LinearRegression ⓘ ?
```

LinearRegression()

8. Get the intercept and the coefficients of the model

```
In [14]: print("Intercept:", model.intercept_)
print("Coefficients:", model.coef_)
```

```
Intercept: -26.27626898612006  
Coefficients: [ 3.03026185e-01  2.19061719e-03 -5.79309274e-02  3.23577722e-01  
   1.66894340e-05 -2.33109315e-03 -3.05992790e-01 -3.15074964e-01]
```

9. Make Predictions using the test data

```
In [15]: # Make predictions on the test set  
# Make predictions on the test set  
y_pred = model.predict(X_test)
```

9. Compare Predictions with Actual Values

```
In [16]: # Create a DataFrame including the test data, and "Actual Price" and "Predicted Price"  
# name the new created DataFrame "results_df"  
  
# Create a DataFrame from X_test  
results_df = X_test.copy()  
  
# Add the actual price values to the new DataFrame results_df into a new column  
results_df['Actual Price'] = y_test  
  
# Add the predicted price values to the new DataFrame results_df into a new column  
results_df['Predicted Price'] = y_pred  
  
# sort results (and display data frame)  
results_df.sort_index()
```

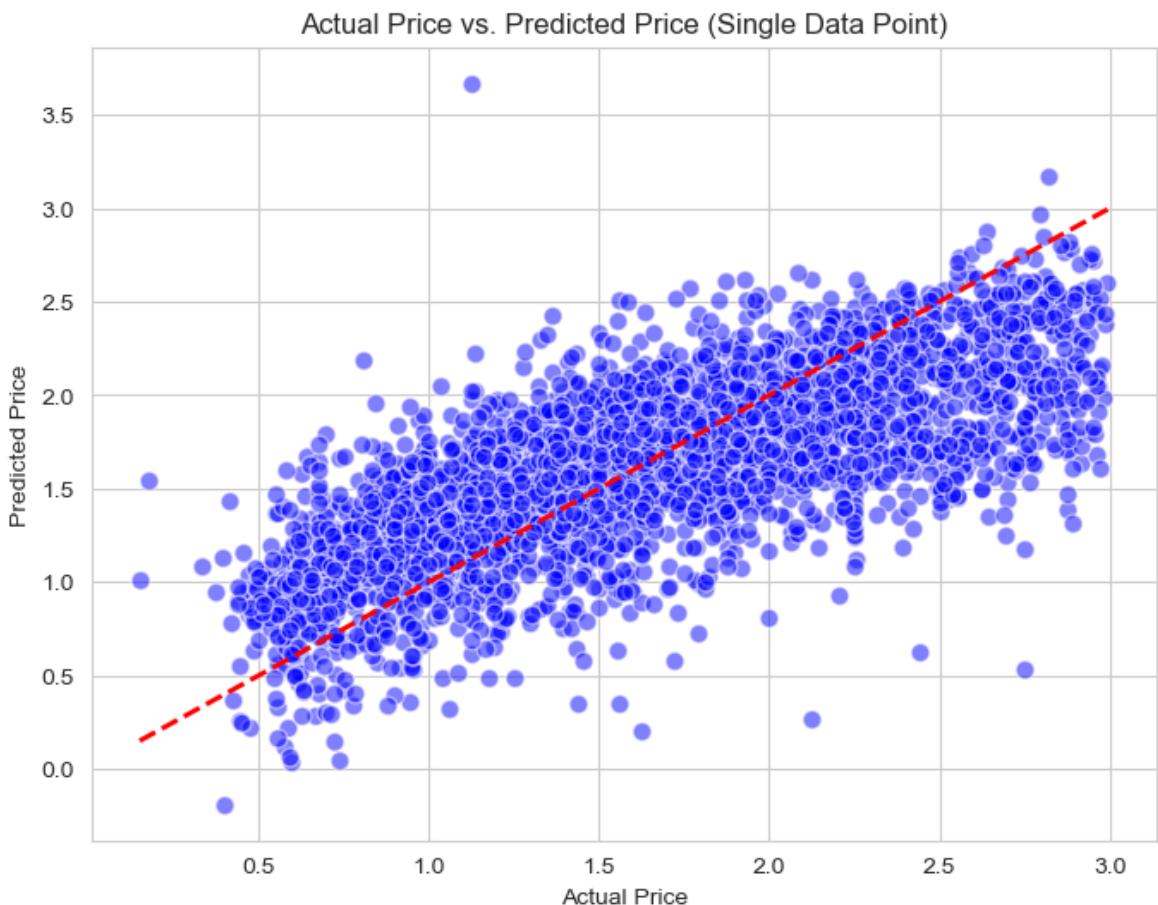
```
Out[16]:
```

	Median Income	Age of House	Average Rooms	Average Bedrooms	Population	Average Occupancy	Latitude	Longitude
5	4.0368	52.0	4.761658	1.103627	413.0	2.139896	37.85	-122.
8	2.0804	42.0	4.294118	1.117647	1206.0	2.026891	37.84	-122.
13	2.6736	52.0	4.000000	1.097701	345.0	1.982759	37.84	-122.
19	2.6033	52.0	5.465455	1.083636	690.0	2.509091	37.84	-122.
24	2.6000	52.0	5.270142	1.035545	1006.0	2.383886	37.84	-122.
...
20616	2.0469	15.0	4.826667	1.176000	1157.0	3.085333	39.08	-121.
20617	3.3021	20.0	4.921053	0.956140	308.0	2.701754	39.06	-121.
20630	3.5673	11.0	5.932584	1.134831	1257.0	2.824719	39.29	-121.
20635	1.5603	25.0	5.045455	1.133333	845.0	2.560606	39.48	-121.
20638	1.8672	18.0	5.329513	1.171920	741.0	2.123209	39.43	-121.

3355 rows × 10 columns

10. Visualization of Actual vs Predicted Values

```
In [17]: # Scatter plot: Actual Price vs. Predicted Price
plt.figure(figsize=(8, 6))
plt.scatter(data=results_df, x='Actual Price', y='Predicted Price', s=60, color='blue')
plt.plot([results_df['Actual Price'].min(), results_df['Actual Price'].max()],
         [results_df['Actual Price'].min(), results_df['Actual Price'].max()],
         color='red', linestyle='--', linewidth=2) # 45-degree reference line
plt.title('Actual Price vs. Predicted Price (Single Data Point)')
plt.xlabel('Actual Price')
plt.ylabel('Predicted Price')
plt.show()
```



11. Model Evaluation

```
In [18]: # Evaluate the model by calculating mse and r2
mse = mean_squared_error(y_test, y_pred)
r2 = r2_score(y_test, y_pred)

# Print the model evaluation results (mse and r2)
print(f"Mean Squared Error (MSE): {mse:.2f}")
print(f"R² Score: {r2:.2f}")
```

Mean Squared Error (MSE): 0.19

R² Score: 0.55

In []:

In []: